



18-441/741: Computer Networks Spring Semester, 2020

Project 3: HTTP Pseudo Streaming

Deadline: 11:59pm EST, Apr. 19, 2020

Questions? Email TA

1. HTTP Server

A HTTP server, often known as a web server, provides contents needed for web browsers. These contents could be a HTML file, a plain text file, an image, audio, video, or just a binary file. In a nutshell, our HTTP server will act as a simple file server. It will take a resource identifier (URI), and locate the content for the URI in a directory relative to a content root. Then generate a HTTP response message which contains information about the content, and the content itself.

2. HTTP Protocol

The Hypertext Transfer Protocol is used to facilitate transfer of hypermedia documents (text, contents, links.) The current protocol is based on RFC 2616 (<http://www.w3.org/Protocols/rfc2616/rfc2616.html>). Please familiarize yourself with the following sections. In the documentation, the server you are implementing in this project is referred to as the "origin server."

- HTTP Message / Request / Response format (4,5,6)
- Status Code
 - o 200 OK (10.2.1)
 - o 206 Partial Content (10.2.7)
 - o 404 Not Found (10.4.5)
 - o 500 Internal Server Error (10.5.1)
- Method
 - o GET (9.3)
- Header
 - o Accept-Ranges (14.5)
 - o Connection (14.10)
 - o Content-Length (14.13)
 - o Content-Range (14.16)
 - o Content-Type (14.17)
 - o Date (14.18)

- Last-Modified (14.29)
 - Range (14.35)
- Entity (7)

3. Server Requirements

Your server should be able to understand and properly generate the status codes/methods/headers listed in the previous sections.

- Successfully handle HTTP/1.1 GET method
 - Generate “404 Not Found” response when the content could not be found
 - Generate “200 OK” when the server can locate the content
 - Generate “206 Partial content” when the request includes Range header (14.35), even if the range starts from 0.
- Make sure that all responses (except status 500) include the Date header. All timestamp in the message must be represented in GMT time zone (RFC 1123-format; 24-hours; abbreviate day of week; month; 4-digit year) For example: Sun, 06 Nov 1994 08:49:13 GMT
- You could also interpret/generate additional headers that may help improve the performance (Cache-Control, Age, If-*, X-Content-Duration) However, they are not required.
- Your server should keep the connection alive after receiving a request (unless closed by the client, in which case, the server must also properly close the connection.) If you need to terminate the connection (due to error or capacity problem), you need to include Connection: close header in the response.
- For benchmarking / testing purpose, always include Connection: Keep-Alive in the response header when the connection remains open after serving a request.
- We will only test your server against a simple binary entity body (identity coding, no entity-header, no chunk encoding – section 4.4 in RFC2616.) This means that the entity length must be equal to the length specified in Content-Length response header. Note that the end of the entity body also indicates the end of the HTTP message in a persistent connection.
- Please also make sure that your server generates/recognizes proper end-of-lines with Carriage Return & Line feed characters (\r\n – ASCII 13 and 10.) Note the extra CR-LF after the request/response header section (just before the entity body).
- For content-type, make sure that your server could support the following media types (based on file extensions or Linux file command) – please refer to IANA for official list of media types (<http://www.iana.org/assignments/media-types/>)
 - text/plain (.txt)
 - text/css (.css)
 - text/html (.htm, .html)
 - image/gif (.gif)
 - image/jpeg (.jpg, .jpeg)
 - image/png (.png)
 - application/javascript (.js)

- application/octet-stream (others)
- video/mp4 (.mp4)
- p video/webm (.webm, .ogg)

4. Content Location

By default, the root directory of all contents (often known as the content root) is located in the directory “**content**” *relative* to the directory from which the server executable file was invoked. For example, if your server is executed from the project directory as followed:

```
/.../project1 $> vodserver #port
```

Then the URL `http://<server-ip-address>:<server-port>/video/video.webm` should correspond to,

```
/.../project1/content/video/video.webm
```

5. Browser Communication

Although implementing a client is not part of the first project. You need to understand clients’ behaviors in order to successfully implement the server part which allow actual HTML 5 video playback on your favorite browser.

Browser Test: We will use Firefox 9+ browser to grade the correctness of your video playback, as follows.

- Firefox will initiate the connection and send a HTTP GET request with URI of a video file. The server should respond with a 200 OK status and start sending the file. However, the browser will preemptively disconnect after received sufficient amount of video header.
- The browser will seek to the end of file to identify the length of the video (by requesting near-end range of the video file)
- Subsequent connection will request a different range of media, starting from the byte which was not buffered.
- In an event of a video seek, the browser will incrementally request further and further range of the file until the frame match the seek time.
- Make sure that you properly handle the Last-Modified response header; otherwise browsers will not be able to cache your video.
- If you are using Chrome, the browser will use multiple connections for a single video transfer. Your server should be able to easily handle concurrent connections.

6. Discussions

I’m confused, where do I start?

Relax; the project may not be as hard (or as easy) as you think initially. Make sure you start early. If you are not familiar with network programming, please consult the materials from the lecture and the RFCs. Once you have thought about your server design, you can start by implementing a core server (the echo

server is a good start) then implement a message parser module which can parse simple request method (just a simple GET request, ignoring other headers). After you have a basic communication server, you could implement the required specification later. Make sure that your server can serve simple file such as html or image files. Then, once you have implemented the server range request, you should be able to test your server against any Firefox 9+ browsers. We will post sample video files for you to test on your own computer.

How do I measure my server performance?

After you are finished with the functionality, you could focus on the performance part. Since your server is standard-compliant, you could use existing tools such as Apache benchmark; ab (<http://httpd.apache.org/docs/2.0/programs/ab.html>) to test the scalability of your server. Note that since ab only simulates a HTTP/1.0 client, you should include `Connection: Keep-Alive` in the response header to prevent it from hanging. Try to see how many concurrent requests (for a 10KB content) you can have, while still allowing 95% of them to be served within 500ms.

Help! I accidentally delete my working code.

Make a lot of backups, if you do not have the source when the deadline comes, you do not have the project. You may also want to use version control tools to manage your source. The ECE undergraduate cluster machines are equipped with both git and subversion tools (although learning to use them is not within the scope of this course). Make sure you have a working branch/revision available for final submission before start tuning for performance or adding features. You could always submit an early working version to the hand-in directory. So make sure you contact the TA and give us information about your team early.

7. Hand-in Procedure

You will submit your project to blackboard and we will test your codes using our own server. Create a file named README in the primary submission directory so we know that it is the primary directory that we need to grade (we will ignore the other).

5. Deliverable Items

The deliverables are enumerated as follows,

1. **The source code** for the entire project (if you use external libraries, provide the binaries needed for compilation of your project)
2. **Makefile** - You should prepare a makefile for C make (or build.xml for Java Ant) which generates the executable file for the project. We expect that the file to work from your submission directory. For example, the following calls should generate an executable vodserver (for C) or VodServer.class (in edu.cmu.ece package directory for Java) respectively. The following commands will be attempted on your submission directory, depending on your choice of programming language.

```
.../project1$ make
.../project1$ ant
```

3. **Do not submit the executable files** (we will DELETE them and deduct your logistic points). However, we must be able to invoke the one of the following commands (after invoking make/ant) from your submission directory to start your server on the given port number (e.g. 8345 in this case).


```
$ vodserver 8345
```

```
$ java edu.cmu.ece.VodServer 8345
```
4. **A brief design document** regarding your server design in design.txt / design.pdf in the submission directory (1-2 pages). Apart from the team information, tell us about the following,
 - a. Your server design, model and components (add pictures if it helps)
 - b. How did you handle multiple clients? How many clients do you think your server can handle effectively?
 - c. Libraries used (optional; name, version, homepage-URL; if available)
 - d. Extra capabilities you implemented (optional, anything which was not specified in the Server Requirement section)
 - e. Extra instructions on how to execute the code, if needed (WARNING: if we cannot compile and run your server on our cluster machine with the above execution method, your logistic points will be deducted. However, please do tell us any runtime parameter your server may have which may help in tuning)

6. Grading

Partial credit will be available according to the following grading scheme,

Items	Points (100 – 441 + 110 – 741)
Logistics	
- Successful submission & compilation	10
- Design documents	10
Server Correctness	
- Handle standard HTTP GET	10
- Handle HTTP error (404)	10
- Handle HTTP Byte-range request	20
- Play video from browser	10
- Allowing video seek from browser	10
Server Performance	
- < 10% CPU utilization when idle	5
- Handle 10 clients simultaneously	15
Required for 18-741 only (bonus for 18-441)	
- Handle 5000 concurrent clients!	10